

In [49]:

```

1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 from sklearn.model_selection import train_test_split
6 from sklearn.linear_model import LinearRegression
7 from sklearn import metrics
8 from sklearn import preprocessing, svm
9

```

In [50]:

```

1 df=pd.read_csv(r"C:\Users\P. VIJAY KUMAR\Downloads\Advertising.csv")
2 df

```

Out[50]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
...
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

200 rows × 4 columns

In [51]:

```

1 df.head()

```

Out[51]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

In [52]: 1 df.tail()

Out[52]:

	TV	Radio	Newspaper	Sales
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

In [53]: 1 df.describe()

Out[53]:

	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000
mean	147.042500	23.264000	30.554000	15.130500
std	85.854236	14.846809	21.778621	5.283892
min	0.700000	0.000000	0.300000	1.600000
25%	74.375000	9.975000	12.750000	11.000000
50%	149.750000	22.900000	25.750000	16.000000
75%	218.825000	36.525000	45.100000	19.050000
max	296.400000	49.600000	114.000000	27.000000

In [54]: 1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0    TV          200 non-null    float64
1    Radio       200 non-null    float64
2    Newspaper   200 non-null    float64
3    Sales       200 non-null    float64
dtypes: float64(4)
memory usage: 6.4 KB
```

In [55]: 1 df.shape

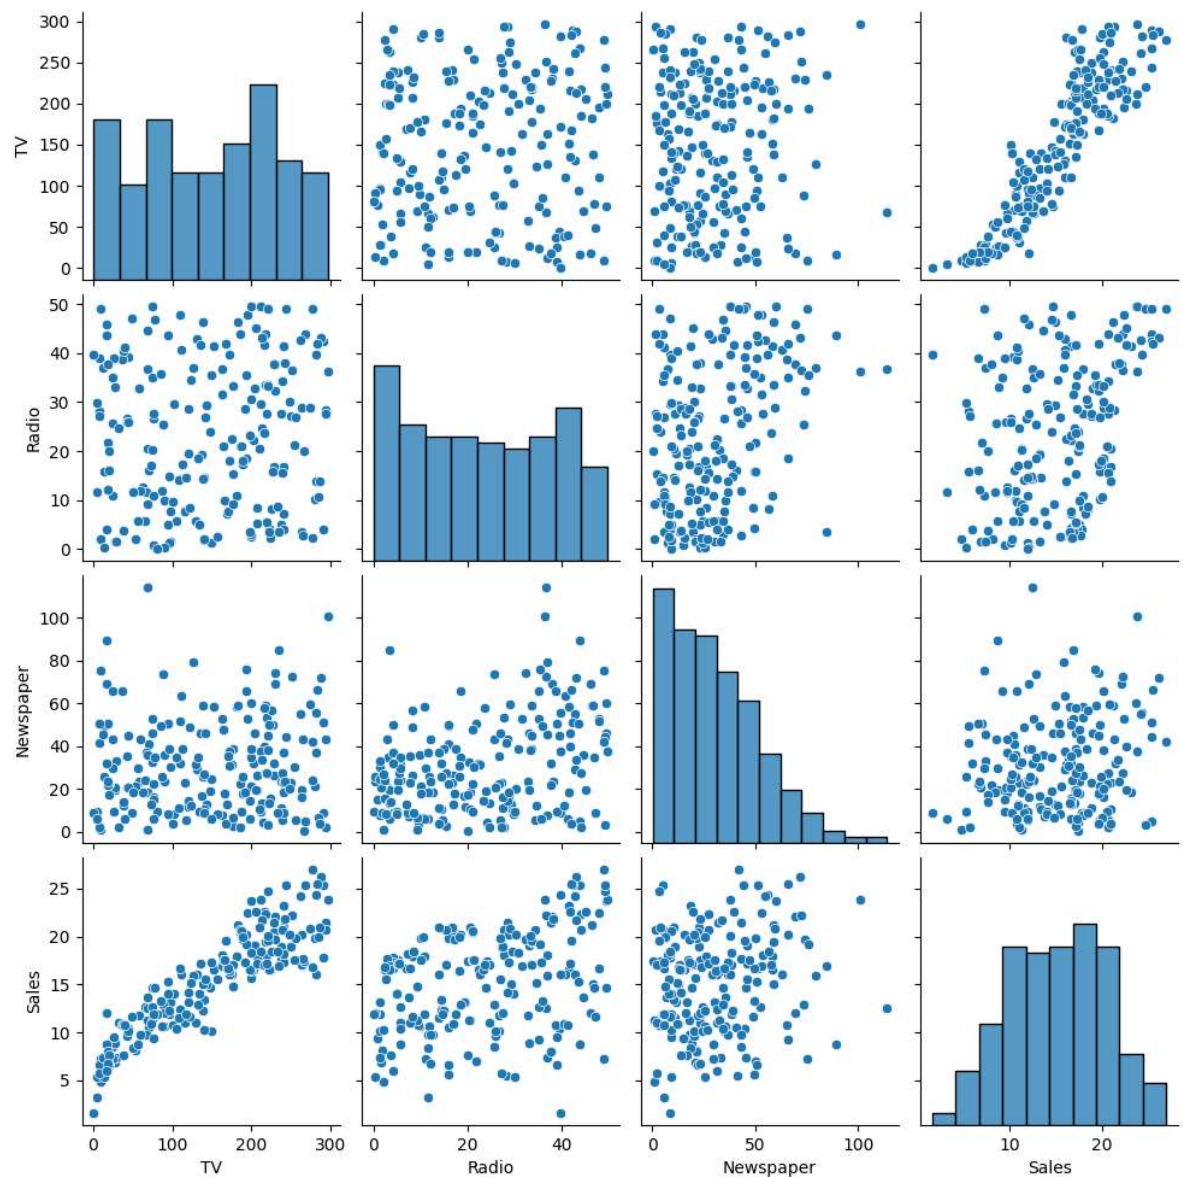
Out[55]: (200, 4)

In [56]: 1 df.columns

Out[56]: Index(['TV', 'Radio', 'Newspaper', 'Sales'], dtype='object')


```
In [57]: 1 sns.pairplot(df)
```

```
Out[57]: <seaborn.axisgrid.PairGrid at 0x212af247910>
```




```
In [58]: 1 sns.displot(df['Sales'])
```

```
Out[58]: <seaborn.axisgrid.FacetGrid at 0x212af6dfb20>
```



```
In [59]: 1 addf=df[['TV', 'Radio', 'Newspaper', 'Sales']]
          2 sns.heatmap(addf.corr())
          3
```

Out[59]: <Axes: >



```
In [60]: 1 X=addf[['TV', 'Radio', 'Newspaper']]
          2 y=df['Sales']
```

```
In [61]: 1 from sklearn.model_selection import train_test_split
          2 X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_s
          3 from sklearn.linear_model import LinearRegression
          4 lm=LinearRegression()
          5 lm.fit(X_train,y_train)
          6 print(lm.intercept_)
```

4.681232151484295

```
In [62]: 1 coeff_df=pd.DataFrame(lm.coef_,X.columns,columns=['coefficient'])
          2 coeff_df
          3
```


Out[62]:

	coefficient
TV	0.054930
Radio	0.109558
Newspaper	-0.006194

```
In [63]: 1 predictions=lm.predict(X_test)
```


```
In [64]: 1 plt.scatter(y_test,predictions)
```

```
Out[64]: <matplotlib.collections.PathCollection at 0x212ae8a2b60>
```



```
In [65]: 1 sns.displot((y_test,predictions),bins=50)#without semicolon
```

```
Out[65]: <seaborn.axisgrid.FacetGrid at 0x212af639ed0>
```



```
In [66]: 1 from sklearn import metrics
2 print('MAE:',metrics.mean_absolute_error(y_test,predictions))
3 print('MSE:',metrics.mean_squared_error(y_test,predictions))
4 print('MAE:',np.sqrt(metrics.mean_squared_error(y_test,predictions)))
```


MAE: 1.3731200698367851

MSE: 2.8685706338964967

MAE: 1.6936855180040056


```
In [67]: 1 sns.lmplot(x="Newspaper",y="Sales",data=df,order=2)
```

```
Out[67]: <seaborn.axisgrid.FacetGrid at 0x212b01c2260>
```





```
In [68]: 1 sns.lmplot(x="TV",y="Sales",data=df,order=2)
```

```
Out[68]: <seaborn.axisgrid.FacetGrid at 0x212b02a87c0>
```



```
In [69]: 1 sns.lmplot(x="Radio",y="Sales",data=df,order=2)
        2
```

```
Out[69]: <seaborn.axisgrid.FacetGrid at 0x212b039e050>
```



```
In [70]: 1 df.fillna(method='ffill',inplace=True)
```

```
In [71]: 1 df.fillna(method='ffill',inplace=True)
```

```
In [72]: 1 regr=LinearRegression()
        2
```

```
In [73]: 1 x=np.array(df['TV']).reshape(-1,1)
        2 y=np.array(df['Sales']).reshape(-1,1)
        3 df.dropna(inplace=True)
```


```
In [74]: 1 X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
2         regr.fit(X_train,y_train)
3         regr.fit(X_train,y_train)
4
```

Out[74]: LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.


On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [75]: 1 y_pred=regr.predict(X_test)
2         plt.scatter(X_test,y_test,color='y')
3         plt.plot(X_test,y_pred,color='b')
4         plt.show()
```



```
In [76]: 1 sns.pairplot(df,x_vars=['TV', 'Radio', 'Newspaper'],y_vars='Sales',height:
```

```
Out[76]: <seaborn.axisgrid.PairGrid at 0x212b04876d0>
```




```
In [77]: 1 #accuracy
2 regr=LinearRegression()
3 regr.fit(X_train,y_train)
4 regr.fit(X_train,y_train)
5 print(regr.score(X_test,y_test))
6
```

```
0.797139039213951
```

```
In [78]: 1 from sklearn.linear_model import Lasso,Ridge
2 from sklearn.preprocessing import StandardScaler
3
```

```
In [79]: 1 ddf=df[['TV', 'Radio', 'Newspaper', 'Sales']]
```

```
In [80]: 1 df.drop(columns = ["Radio", "Newspaper"], inplace = True)
2 sns.pairplot(df)
3 df.Sales=np.log(df.Sales)
4
```



```
In [81]: 1 features=df.columns[0:2]
2 target=df.columns[-1]
3 X=df[features].values
4 y=df[target].values
5 X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_s
6 print("The dimension of X_train is {}".format(X_train.shape))
7 print("The dimension of X_test is {}".format(X_test.shape))
8 scaler=StandardScaler()
9 X_train=scaler.fit_transform(X_train)
10 X_test=scaler.transform(X_test)
```

The dimension of X_train is (140, 2)

The dimension of X_test is (60, 2)

```
In [82]: 1 #Linear regression model
2 regr=LinearRegression()
3 regr.fit(X_train,y_train)
4 actual=y_test #actual value
5 train_score_regr=regr.score(X_train,y_train)
6 test_score_regr=regr.score(X_test,y_test)
7 print("\nLinear model:\n")
8 print("The train score for Linear model is {}".format(train_score_regr))
9 print("The test score for Linear model is {}".format(test_score_regr))
10
```

Linear model:

The train score for Linear model is 1.0

The test score for Linear model is 1.0

```
In [83]: 1 #ridge regression model
2 ridgeReg=Ridge(alpha=10)
3 ridgeReg.fit(X_train,y_train)
4 #train and test score for ridge regression
5 train_score_ridge=ridgeReg.score(X_train,y_train)
6 test_score_ridge=ridgeReg.score(X_test,y_test)
7 print("\nRidge model:\n")
8 print("The train score for ridge model is {}".format(train_score_ridge))
9 print("The test score for ridge model is {}".format(test_score_ridge))
10
```

Ridge model:

The train score for ridge model is 0.9902871391941609

The test score for ridge model is 0.984426628514122

```
In [84]: 1 #using the linear cv model for ridge regression
2 from sklearn.linear_model import RidgeCV
3 #ridge cross validation
4 ridge_cv=RidgeCV(alphas=[0.0001,0.001,0.01,0.1,1,10]).fit(X_train,y_train)
5 #score
6 print(ridge_cv.score(X_train,y_train))
7 print(ridge_cv.score(X_test,y_test))
```

0.999999999997627


0.9999999999962466

```
In [85]: 1 #using the linear cv model for lasso regression
2 from sklearn.linear_model import LassoCV
3 #lasso cross validation
4 lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,0.1,1,10],random_state=0).fit(X_train,y_train)
5 #score
6 print(lasso_cv.score(X_train,y_train))
7 print(lasso_cv.score(X_test,y_test))
8
```

0.9999999343798134


0.9999999152638072

```
In [87]: 1 plt.figure(figsize=(10,10))
2 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',mar
3          label=r'Ridge;\alpha=10$',zorder=7)
4 plt.plot(features,regr.coef_,alpha=0.5,linestyle='none',marker='*',marker:
5          color='green',label='Linear Regression')
6 plt.xticks(rotation=90)
7 plt.legend()
8 plt.show()
```




```
In [88]: 1 #ridge regression  
2 plt.figure(figsize=(10,10))  
3 sns.heatmap(ddf.corr(),annot=True)
```

Out[88]: <Axes: >




```
In [89]: 1 #Lasso regression model
2 lassoReg=Lasso(alpha=10)
3 lassoReg.fit(X_train,y_train)
4 #train and test score for ridge regression
5 train_score_lasso=lassoReg.score(X_train,y_train)
6 test_score_lasso=lassoReg.score(X_test,y_test)
7 print("\nLasso model:\n")
8 print("The train score for lasso model is {}".format(train_score_lasso))
9 print("The test score for lasso model is {}".format(test_score_lasso))
```

Lasso model:

The train score for lasso model is 0.0


The test score for lasso model is -0.0042092253233847465

```
In [90]: 1 plt.figure(figsize=(10,10))
2 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',mar
3          label=r'Ridge;\alpha=10$',zorder=7)
4 plt.plot(features,regr.coef_,alpha=0.5,linestyle='none',marker='*',marker:
5          color='green',label='Linear Regression')
6 plt.xticks(rotation=90)
7 plt.legend()
8 plt.show()
```



```
In [91]: 1 pd.Series(lassoReg.coef_,features).sort_values(ascending=True).plot(kind=
```

```
Out[91]: <Axes: >
```




In [92]:

```

1  #plot size
2  plt.figure(figsize=(10,10))
3  #add plot for ridge regression
4  plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',mar
5           color='red',label=r'Ridge;$\alpha=10$',zorder=7)
6  #add plot for lasso regression
7  plt.plot(features,lassoReg.coef_,alpha=0.7,linestyle='none',marker='*',mar
8           color='green',label=r'Lasso;$\alpha=10$',zorder=7)
9  #add plot for linear model
10 plt.plot(features,regr.coef_,alpha=0.5,linestyle='none',marker='*',marker:
11           color='b',label=r'LinearRegression')
12 #rotate axis
13 plt.xticks(rotation=90)
14 plt.legend()
15 plt.title("Comparison of Ridge,Lasso and Linear regression models")
16 plt.show()

```



```
In [95]: 1 #Elasticnet
2 from sklearn.linear_model import ElasticNet
3 regr=ElasticNet()
4 regr.fit(X,y)
5 print(regr.coef_)
6 print(regr.intercept_)
7 y_pred_elastic=regr.predict(X_train)
8 mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
9 print("Mean Squared Error on test set",mean_squared_error)
```

```
[0.00417976 0.          ]
2.026383919311004
Mean Squared Error on test set 0.5538818050142158
```

```
In [ ]: 1
```